# LG-TargetSeg: Deep Semantic Fusion with Language-Guided State-Space Modeling for Targeted Medical Image Segmentation – Supplementary Material

Mohamed Lamine Allaoui and Mohand Saïd Allili

## A Universal Training datasets and experimental setup

### A.1 Complete Dataset Specifications for Universal Generalization

Evaluating the generalizability of cross-modal fusion models across diverse clinical domains requires training data that spans multiple imaging modalities, anatomical systems, and pathological conditions. While task-specific evaluation provides insights under controlled conditions, clinical reality demands robust adaptation across unseen modalities and anatomies. To this end, we developed a comprehensive universal training framework that leverages heterogeneous medical imaging datasets to evaluate LG-TargetSeg's cross-domain fusion and generalization capacity without architectural retraining.

Our universal training corpus represents a strategically curated collection spanning the breadth of clinical practice. Rather than simply aggregating available datasets, we selected collections that collectively represent major imaging modalities, anatomical systems, and pathological conditions encountered in modern radiology. This principled approach ensures our generalization evaluation reflects realistic clinical deployment scenarios while maintaining the semantic flexibility essential for prompt-driven cross-modal fusion.

Table 1: Universal Training Dataset Specifications

| **Dataset** | **Clinical Description** | **Studies** | **Modality** | **Primary Anatomical Focus** |
|---|---|---|---|---|
| AbdomenCT-1K | Multi-organ abdominal segmentation covering hepatic, renal, and splenic structures | 361 | CT | Liver, kidneys, spleen, pancreas, gallbladder |
| AMOS | Cross-modal abdominal organ segmentation for modality-invariant feature learning | 240 | CT, MRI | Abdominal organs |
| BraTS | Brain tumor segmentation including multiple tumor sub-regions and peritumoral structures | 6,096 | Multi-modal MRI | Gliomas, enhancing regions, tumor core, edema |
| CHAOS | Cross-modal abdominal organ challenge across imaging protocols | 40 | CT, T2-MRI | Liver, kidneys, spleen |
| KiTS 2023 | Kidney and renal tumor segmentation with pathological annotations | 489 | CT | Kidneys, renal tumors, cystic lesions |
| LiTS | Liver tumor segmentation with boundary delineation in oncological imaging | 131 | CT | Liver parenchyma, hepatic lesions |
| LUNA | Pulmonary nodule analysis with lung parenchyma coverage | 888 | CT | Lung tissue, pulmonary nodules |
| MSD Collection | Medical Segmentation Decathlon spanning ten anatomical tasks | 3,225 | CT, MRI | Heart, hippocampus, liver, lung, pancreas, prostate, colon, vessels, spleen |
| PROMISE12 | Prostate segmentation using T2-weighted MRI | 37 | T2-MRI | Prostate gland, peripheral zone |
| ABUS | Automated breast ultrasound with 3D tumor characterization | 200 | 3D US | Breast lesions, fibrocystic changes |
| BUSI | Breast ultrasound with benign and malignant lesion representation | 1,312 | 2D US | Breast masses, tumor boundaries |
| TotalSegmentator V2 | Anatomical atlas segmentation with 117 distinct structures | 1,228 | CT | Full-body anatomical mapping |
| CAMUS | Cardiac motion analysis using multi-view echocardiography | 500 | US | Left ventricle, left atrium, myocardium |
| autoPET 2022 | Whole-body lesion segmentation in oncological PET | 1,014 | PET/CT | Metabolically active lesions, tumor burden |

The foundation of our universal evaluation rests on fourteen distinct medical imaging collections, encompassing over 15,761 individual studies across three primary modalities and providing unprecedented scope for cross-domain learning evaluation. Table 1 presents complete specifications of our universal training corpus, highlighting the diversity of anatomical targets and imaging protocols represented.

This heterogeneous composition serves multiple scientific purposes beyond simple modality variation. It enables evaluation of LG-TargetSeg's ability to fuse linguistic, visual, and domain-specific information into modality-invariant anatomical representations that generalize across imaging physics, from high-resolution CT cross-sectional anatomy to MRI soft tissue characterization, ultrasound real-time acoustic imaging, and PET functional metabolic assessment. The diverse anatomical coverage tests whether the continuous cross-modal fusion mechanism can transfer learned features between different organ systems and pathological conditions, while variation in dataset sizes and annotation quality reflects realistic clinical data collection constraints.

The anatomical distribution within our universal corpus reflects clinical imaging prevalence while ensuring adequate representation across major organ systems. Neurological studies comprise the largest component at 40.6% of the corpus, primarily driven by the comprehensive BraTS brain tumor dataset, encompassing both structural and pathological brain imaging. Multi-organ anatomical atlas studies represent 24.4%, providing comprehensive full-body coverage through TotalSegmentator and MSD collections. Thoracic studies account for 10.7%, including both pulmonary and cardiac applications. Breast imaging represents 9.6% through ultrasound modalities, while abdominal imaging comprises 8.0% covering hepatic, renal, and splenic structures. Oncological lesion studies contribute 6.4% through whole-body PET/CT imaging. This distribution reflects the comprehensive scope of our universal training while providing sufficient diversity for robust cross-domain evaluation.

## A.2 Universal Data Preparation and Text Prompt Generation

Preparing medical imaging data for universal multi-source fusion training presents unique challenges extending beyond conventional computer vision preprocessing. Medical images exhibit profound variations in intensity ranges,

spatial resolutions, and acquisition parameters reflecting underlying imaging physics, while clinical interpretation relies on domain-specific knowledge that must be preserved throughout preprocessing. Our approach addresses these challenges through a comprehensive framework that normalizes technical variations while preserving the clinically relevant information that the fusion mechanism depends on.

### A.2.1 Strategic Slice Extraction and Multi-Label Task Construction

Given LG-TargetSeg's two-dimensional architecture, we extract representative slices from three-dimensional volumes using a dual-strategy approach balancing anatomical coverage with label representation. We extract two distinct slice types: mid-slices and max-slices, each serving specific training purposes.

Mid-slice extraction systematically captures the middle slice along each anatomical axis (axial, sagittal, coronal), resulting in representative $256 \times 256$ cross-sections. This strategy avoids bias toward pathological regions, ensuring LG-TargetSeg learns robust anatomical representations that generalize across spatial locations, critical for clinical deployment where structure locations are unknown *a priori*.

Max-slice extraction supplements training data by identifying slices containing maximum voxel counts for each anatomical label $l$. This addresses label sparsity common in medical imaging, where structures may occupy minimal volume fractions, providing concentrated anatomical examples while maintaining context.

### A.2.2 Combinatorial Prompt Generation and Multi-Class Task Decomposition

A critical component of our multi-source fusion pipeline lies in systematically generating anatomically grounded text prompts that reflect complex multi-class medical annotations. Rather than treating multi-class segmentation as single tasks, we decompose annotated slices into multiple prompt-driven fusion tasks capturing individual structures and clinically relevant combinations.

For datasets with multiple anatomical labels per slice, we implement combinatorial task generation creating individual and composite targets. The

CAMUS cardiac dataset exemplifies this approach: containing left ventricular epicardium, endocardium, and left atrium, we generate binary tasks reflecting clinical reasoning patterns rather than single three-class segmentation.

Individual structures receive dedicated binary masks with corresponding prompts ("Segment the left ventricular epicardium," "Delineate the left atrial boundary"), enabling precise anatomical discrimination with clinical terminology consistency. Combinatorial tasks reflect meaningful anatomical relationships through composite prompts ("Segment the epicardium and endocardium"), with binary masks encompassing specified region unions.

This approach serves multiple purposes: significantly expanding the training corpus, enabling the fusion mechanism to learn flexible semantic associations across anatomical specificity levels, and mirroring clinical workflows where radiologists focus on individual or combined structures based on clinical questions.

Mathematically, for slices containing $n$ anatomical labels $\{l_1, l_2, \ldots, l_n\}$, we generate $2^n - 1$ distinct tasks corresponding to all non-empty label subsets. Each task $t_S$ for subset $S$ receives binary mask $M_S = \bigcup_{l_i \in S} M_{l_i}$ and corresponding textual prompt $P_S$ describing the anatomical union in clinically-appropriate language.

# B Reinforced Gaussian Dice Loss (RGDL) Extended Analysis

## B.1 Mathematical Formulation for Binary Segmentation

### B.1.1 Notation and Preliminaries

We establish the following mathematical notation for the segmentation task:

- $X \in \mathbb{R}^{B \times C \times H \times W}$: Batch of input images, where $B$ denotes the batch size, $C$ represents the number of channels, and $H, W$ correspond to the height and width dimensions
- $Y \in \{0, 1\}^{B \times 1 \times H \times W}$: Original binary ground truth segmentation masks
- $\tilde{Y} \in [0, 1]^{B \times 1 \times H \times W}$: Gaussian-smoothed ground truth masks
- $\hat{Y} \in [0, 1]^{B \times 1 \times H \times W}$: Predicted probability maps
- $f_\theta : X \to \hat{Y}$: Segmentation model parameterized by $\theta$

### B.1.2 Gaussian Smoothing of Ground Truth

To account for annotation uncertainty, we apply Gaussian smoothing to the binary ground truth masks:

$$\tilde{Y} = Y * G_\sigma \tag{1}$$

where $*$ represents the convolution operation and $G_\sigma$ is a Gaussian kernel with standard deviation $\sigma$:

$$G_\sigma(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}} \tag{2}$$

For a specific pixel location $(i, j)$, the smoothed ground truth is computed as:

$$\tilde{Y}(i, j) = \sum_{m=-k}^{k} \sum_{n=-k}^{k} Y(i + m, j + n) \cdot G_\sigma(m, n) \tag{3}$$

with $k$ representing the kernel radius (typically $k = 3\sigma$).

This smoothing operation converts hard binary masks to soft probability masks, with values ranging from 0 to 1, representing the confidence or probability that a pixel belongs to the foreground. Pixels at object boundaries will have intermediate values, reflecting the uncertainty in the exact boundary location, resulting in a dilated mask.Figure 1 illustrates the effect of different sigma values on mask dilation, demonstrating how Gaussian smoothing progressively expands the segmentation boundary as sigma increases.

### B.1.3 Reinforced Gaussian Dice Loss Formulation

The Reinforced Gaussian Dice Loss is formulated as a weighted combination of two complementary components:

$$\mathcal{L}_{\text{RGDL}}(\tilde{Y}, Y, \hat{Y}) \;=\; \lambda_1 \cdot \mathcal{L}_{\text{Dice}}(\tilde{Y}, Y, \hat{Y}) \;+\; \lambda_2 \cdot \mathcal{L}_{\text{RW}}(Y, \hat{Y}) \tag{4}$$

where $\lambda_1, \lambda_2 \in [0, 1]$ are weighting coefficients with the constraint $\lambda_1 + \lambda_2 = 1$, enabling adaptive balancing between global and local optimization objectives.

Figure 1: Visualization of Gaussian-based mask dilation at $256 \times 256$ resolution. From left to right: original binary mask, followed by dilated masks with increasing sigma values ($\sigma = 3, 10, 20, 30$). Orange regions show the probability field generated by Gaussian smoothing, with lighter to darker orange indicating lower to higher probability values. Red contours mark the boundary of the dilated mask at 0.5 threshold. Note how larger sigma values produce more extensive boundary expansion while preserving the overall shape characteristics of the original segmentation.

**Modified Dice Loss Component.** We modify the Dice loss term to use the smoothed ground truth only in the numerator while retaining the original binary mask in the denominator:

$$\mathcal{L}_{\text{Dice}}(\tilde{Y}, Y, \hat{Y}) \;=\; 1 \;-\; \frac{2\sum_{i=1}^{N} \tilde{y}_i \hat{y}_i + \epsilon}{\sum_{i=1}^{N} y_i \;+\; \sum_{i=1}^{N} \hat{y}_i \;+\; \epsilon}, \tag{5}$$

where $\tilde{y}_i$, $y_i$, and $\hat{y}_i$ represent individual elements of $\tilde{Y}$, $Y$, and $\hat{Y}$ respectively, $\epsilon$ constitutes a small smoothing constant to ensure numerical stability by preventing division by zero, and $N = B \times H \times W$ denotes the total number of pixels across the batch.

This formulation offers a significant advantage: it encourages the model to extend predicted regions slightly beyond the hard boundaries (improving recall) while still maintaining volumetric consistency with the original mask. By using the original binary mask in the denominator, we ensure that the total volume of the prediction remains comparable to the ground truth, preventing excessive over-segmentation that might occur if the smoothed mask were used in both numerator and denominator.

**Learnable Reward-Weighting Parameters.** To enhance the adaptability of our loss function, we replace the fixed hyperparameters with learnable parameters. We define:

$$\alpha = \alpha_{\max} \cdot \sigma(\theta_\alpha) \tag{6}$$
$$\beta = \beta_{\max} \cdot \sigma(\theta_\beta) \tag{7}$$
$$\gamma = \gamma_{\max} \cdot \sigma(\theta_\gamma) \tag{8}$$

where $\sigma(x) = \frac{1}{1+e^{-x}}$ is the sigmoid function, ensuring that the parameters remain within a reasonable range. The terms $\alpha_{\max}$, $\beta_{\max}$, and $\gamma_{\max}$ are scaling constants that define the maximum values these parameters can reach. The parameters $\theta_\alpha$, $\theta_\beta$, and $\theta_\gamma$ are trainable network parameters that are learned through backpropagation.

**Reward-Weighted Component.** To systematically incorporate pixel-wise incentives and penalties with learnable parameters, we define a reward function that differentiates between various prediction outcomes:

$$r(y_i, \hat{y}_i) = \begin{cases} +\alpha \cdot y_i \hat{y}_i & \text{(true positive contribution)}, \\ -\beta \cdot (1 - y_i)\hat{y}_i & \text{(false positive contribution)}, \\ -\gamma \cdot y_i(1 - \hat{y}_i) & \text{(false negative contribution)}, \end{cases} \tag{9}$$

We use the original binary ground truth values $y_i$ to maintain clear and targeted reward signals that precisely distinguish between different prediction outcomes.

Based on this reward structure, we define the reward-weighted loss term by aggregating the rewards/penalties and normalizing by the total number of positive (foreground) pixels in the original ground truth:

$$\mathcal{L}_{\text{RW}}(Y, \hat{Y}) \;=\; -\frac{1}{\sum_{i=1}^{N} y_i + \epsilon}\Big(\alpha \sum_{i=1}^{N} y_i \hat{y}_i \;-\; \beta \sum_{i=1}^{N} (1-y_i)\, \hat{y}_i \;-\; \gamma \sum_{i=1}^{N} y_i (1-\hat{y}_i)\Big). \tag{10}$$

The negative sign translates the reward-maximization objective into the standard loss-minimization framework.

**Combined Reinforced Gaussian Dice Loss.** Integrating the modified Dice and reward-weighted components, the complete Reinforced Gaussian

Dice Loss is expressed as:

$$\mathcal{L}_{\text{RGDL}}(\tilde{Y}, Y, \hat{Y}) = \lambda_1 \Big[1 - \frac{2\sum_{i=1}^{N} \tilde{y}_i \hat{y}_i + \epsilon}{\sum_{i=1}^{N} y_i + \sum_{i=1}^{N} \hat{y}_i + \epsilon}\Big] \tag{11}$$

$$- \lambda_2 \frac{\alpha \sum_{i=1}^{N} y_i \hat{y}_i \; - \; \beta \sum_{i=1}^{N} (1 - y_i)\hat{y}_i \; - \; \gamma \sum_{i=1}^{N} y_i (1 - \hat{y}_i)}{\sum_{i=1}^{N} y_i + \epsilon} \tag{12}$$

where $\alpha$, $\beta$, and $\gamma$ are the learnable parameters as defined above. This formulation establishes a sophisticated balance between global structural overlap assessment (via the Dice component) and localized pixel-wise supervisory signals (via the reward-weighted component), while also accounting for annotation uncertainty through Gaussian smoothing.

### B.1.4 Differentiability and Gradients

The Reinforced Gaussian Dice Loss maintains full differentiability with respect to the network parameters $\theta$, as each component exhibits smooth dependence on the predicted probabilities $\hat{y}_i$. For a specific pixel's prediction $\hat{y}_j$, the gradient is computed as:

$$\frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \hat{y}_j} = \lambda_1 \frac{\partial \mathcal{L}_{\text{Dice}}}{\partial \hat{y}_j} + \lambda_2 \frac{\partial \mathcal{L}_{\text{RW}}}{\partial \hat{y}_j}. \tag{13}$$

**Gradient of the Modified Dice Component.** For the modified Dice term $\mathcal{L}_{\text{Dice}} = 1 - \frac{2\sum \tilde{y}_i \hat{y}_i + \epsilon}{\sum y_i + \sum \hat{y}_i + \epsilon}$, application of the quotient rule yields:

$$\frac{\partial \mathcal{L}_{\text{Dice}}}{\partial \hat{y}_j} = - \frac{\partial}{\partial \hat{y}_j} \Big[\frac{2\sum_{i=1}^{N} \tilde{y}_i \hat{y}_i + \epsilon}{\sum_{i=1}^{N} y_i + \sum_{i=1}^{N} \hat{y}_i + \epsilon}\Big] \tag{14}$$

$$= - \frac{2\,\tilde{y}_j \big(\sum y_i + \sum \hat{y}_i + \epsilon\big) \; - \; \big(2\sum_{i=1}^{N} \tilde{y}_i \hat{y}_i + \epsilon\big)}{\big(\sum y_i + \sum \hat{y}_i + \epsilon\big)^2}. \tag{15}$$

This gradient formulation reveals a key advantage: pixels near object boundaries where $\tilde{y}_j > 0$ receive a stronger positive gradient signal than they would with the conventional Dice loss. This specifically encourages the model to expand its predictions slightly beyond the hard boundary defined in the original binary mask, improving recall without compromising overall volumetric consistency.

**Gradient of the RW Component.** For the reward-weighted term with original binary ground truth

$$\mathcal{L}_{\text{RW}} = -\frac{1}{\sum y_i + \epsilon}\Big[\alpha \sum y_i \hat{y}_i - \beta \sum (1-y_i)\hat{y}_i - \gamma \sum y_i(1-\hat{y}_i)\Big],$$

we conduct a case-based analysis depending on whether the pixel belongs to the foreground or background:

**Case 1: $y_j = 1$ (Foreground Pixel).** For foreground pixels, only the terms $\alpha \sum y_i \hat{y}_i$ and $-\gamma \sum y_i(1-\hat{y}_i)$ contribute to the gradient calculation for pixel $j$. The derivative is computed as:

$$\frac{\partial \mathcal{L}_{\text{RW}}}{\partial \hat{y}_j} = -\frac{1}{\sum y_i + \epsilon}\Big[\alpha \frac{\partial}{\partial \hat{y}_j}(y_j \hat{y}_j) + \big(-\gamma\big) \frac{\partial}{\partial \hat{y}_j}\big(y_j(1-\hat{y}_j)\big)\Big]. \tag{16}$$

Given that $y_j = 1$, we have $\frac{\partial}{\partial \hat{y}_j}(y_j \hat{y}_j) = 1$ and $\frac{\partial}{\partial \hat{y}_j}(y_j(1-\hat{y}_j)) = -1$. Consequently, the expression within the brackets simplifies to $\alpha + \gamma$, resulting in:

$$\frac{\partial \mathcal{L}_{\text{RW}}}{\partial \hat{y}_j} = -\frac{\alpha + \gamma}{\sum_{i=1}^{N} y_i + \epsilon} \quad (\text{for } y_j = 1).$$

This formulation reveals an important insight: both increasing true positives (rewarded by $\alpha$) and decreasing false negatives (penalized by $\gamma$) contribute additively to the gradient signal, enhancing the model's ability to identify foreground structures.

**Case 2: $y_j = 0$ (Background Pixel).** For background pixels, only the term $-\beta \sum (1-y_i)\hat{y}_i$ contributes to the gradient calculation for pixel $j$:

$$\frac{\partial \mathcal{L}_{\text{RW}}}{\partial \hat{y}_j} = -\frac{1}{\sum y_i + \epsilon}\Big[-\beta \frac{\partial}{\partial \hat{y}_j}\big((1-y_j)\hat{y}_j\big)\Big].$$

Since $y_j = 0$ implies $(1-y_j) = 1$, the derivative $\frac{\partial}{\partial \hat{y}_j}(\hat{y}_j) = 1$, yielding:

$$\frac{\partial \mathcal{L}_{\text{RW}}}{\partial \hat{y}_j} = +\frac{\beta}{\sum_{i=1}^{N} y_i + \epsilon}.$$

This positive gradient indicates that increasing the prediction value for background pixels is appropriately penalized by factor $\beta$, discouraging false positive predictions.

**Gradients for Learnable Parameters.** The loss function also produces gradients with respect to the learnable reward-weighting parameters:

$$\frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \theta_\alpha} = \frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \alpha} \cdot \frac{\partial \alpha}{\partial \theta_\alpha} \tag{17}$$

$$\frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \theta_\beta} = \frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \beta} \cdot \frac{\partial \beta}{\partial \theta_\beta} \tag{18}$$

$$\frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \theta_\gamma} = \frac{\partial \mathcal{L}_{\text{RGDL}}}{\partial \gamma} \cdot \frac{\partial \gamma}{\partial \theta_\gamma} \tag{19}$$

With the sigmoid parameterization:

$$\frac{\partial \alpha}{\partial \theta_\alpha} = \alpha_{\max} \cdot \sigma(\theta_\alpha) \cdot (1 - \sigma(\theta_\alpha)) \tag{20}$$

$$\frac{\partial \beta}{\partial \theta_\beta} = \beta_{\max} \cdot \sigma(\theta_\beta) \cdot (1 - \sigma(\theta_\beta)) \tag{21}$$

$$\frac{\partial \gamma}{\partial \theta_\gamma} = \gamma_{\max} \cdot \sigma(\theta_\gamma) \cdot (1 - \sigma(\theta_\gamma)) \tag{22}$$

These gradients enable the reward-weighting parameters to be learned through backpropagation along with the network weights.

## B.2 Hyperparameter Configuration

The Reinforced Gaussian Dice Loss requires configuration of several hyperparameters to optimize performance for specific segmentation tasks:

- $\lambda_1, \lambda_2 \in [0, 1]$: These coefficients control the relative weighting between the Dice and reward-weighted components, constrained such that $\lambda_1 + \lambda_2 = 1$ to maintain a balanced optimization objective.
- $\sigma$: This parameter controls the standard deviation of the Gaussian kernel used for smoothing the ground truth masks. Larger values of $\sigma$ create more uncertainty around boundaries, which can help in cases where annotations are particularly variable. In all experiments we set $\sigma = 3$.
- $\alpha_{\max}$: This parameter determines the maximum reward magnitude for true positive predictions. In all experiments we set $\alpha_{\max} = 1.0$.
- $\beta_{\max}$, $\gamma_{\max}$: These coefficients establish the maximum penalty magnitudes for false positive and false negative predictions, respectively. We use $\beta_{\max} = \gamma_{\max} = \alpha_{\max} = 1.0$, allowing the network to learn the asymmetry during training.

- $\epsilon$: This smoothing constant is set to $10^{-5}$ to ensure numerical stability during optimization.

**Parameter Initialization.** The learnable parameters $\theta_\alpha$, $\theta_\beta$, and $\theta_\gamma$ are all initialized to zero. Because the sigmoid function satisfies $\sigma(0) = 0.5$, this yields initial values of $\alpha = \beta = \gamma = 0.5 \cdot \alpha_{\max} = 0.5$, providing a symmetric starting point. During training, the network discovers the relative cost of each error type and the three weights diverge accordingly. In practice, $\gamma$ tends to increase above $\beta$, reflecting the clinical preference for penalizing false negatives (missed pathology) more than false positives (over-segmentation).

For multi-class segmentation tasks, class-specific parameter sets $(\alpha_k, \beta_k, \gamma_k)$ can be employed, each initialized to zero and learned independently per class.

## B.3 Extension to Multi-class Segmentation

For segmentation tasks involving multiple classes, the Reinforced Gaussian Dice Loss can be naturally extended to accommodate K-class problems. Let $Y_k$, $\tilde{Y}_k$ and $\hat{Y}_k$ represent the original binary ground-truth mask, smoothed ground-truth mask, and predicted probability map for class $k$, respectively. The multi-class formulation is defined as:

$$\mathcal{L}_{\text{MC-RGDL}}(\tilde{Y}, Y, \hat{Y}) = \sum_{k=0}^{K-1} w_k \Big[ \lambda_1 \, \mathcal{L}_{\text{Dice}}(\tilde{Y}_k, Y_k, \hat{Y}_k) + \lambda_2 \, \mathcal{L}_{\text{RW}}(Y_k, \hat{Y}_k) \Big], \quad (23)$$

where $w_k \geq 0$ represents the weight for class $k$ with the constraint $\sum_{k=0}^{K-1} w_k = 1$.

This formulation enables class-specific error prioritization by allowing each class to have its own set of learnable reward-weighting parameters $(\alpha_k, \beta_k, \gamma_k)$, which is especially valuable in medical imaging contexts where different anatomical structures or pathological regions may have varying clinical significance.